

gevent, threads & async frameworks

Denis Bilenko

What is gevent?

- green threads for Python
 - Cooperative, only switch at I/O
- greenlet for context switching
 - No syscalls
- libev for event loop
 - Pretty fast

status

- current stable: 1.0.1
 - Runs on Python 2.5 – 2.7
- master branch: 1.1
 - Support for PyPy: 85% tests pass
 - Support for Python3: 53% tests pass
 - Basics and sockets work
 - Subprocesses and fileobjects todo

```
import thread, socket
```

```
def send(host):
```

```
    sock = socket.create_connection((host, 80))
```

```
    sock.sendall("GET / HTTP/1.0\r\n\r\n")
```

```
# do two requests in parallel
```

```
create_new_thread(send, ('python.org', ))
```

```
create_new_thread(send, ('gevent.org', ))
```

```
#import thread, socket
from gevent import thread, socket

def send(host):
    sock = socket.create_connection((host, 80))
    sock.sendall("GET / HTTP/1.0\r\n\r\n")

# do two requests in parallel
create_new_thread(send, ('python.org', ))
create_new_thread(send, ('gevent.org', ))
```

drop-in modules

```
from gevent import  
    socket,  
    ssl,  
    subprocess,  
    thread,  
    local,  
    queue
```

stdlib compatibility

- less to learn
- API is more stable across gevent versions
- we can use stdlib's tests to verify semantics
- trivial to port libraries and apps to gevent

```
from gevent import monkey; monkey.patch_all()

import requests
from flask import Flask
app = Flask(__name__)

@app.route('/')
def hello_world():
    return requests.get('http://python.org').content

app.run()
```


A TCP server

```
def echo(socket, address):  
    for line in socket.makefile():  
        socket.sendall(line)  
  
gevent.server.StreamServer(':5000', handle).start()
```

`gevent.wait([... objects ...])`

- wait for any gevent object
- extendable through `rawlink(callback)`

```
gevent.wait([... objects ...], timeout=5)
```

- limit waiting to certain time

```
gevent.wait([... objects ...], count=N)
```

- only wait for N objects
- return value is a list of ready objects

gevent.wait()

- wait for everything
- “background” watchers `ref=False`
- graceful shutdown:
 - stop accepting new requests
 - `gevent.wait()`

geventserver

- pre-fork, <999LOC
- supports any gevent server, not just HTTP
- graceful shutdown
- can be embedded: `import geventserver`
- reports long loop iterations

github.com/surfly/geventserver

Why?

- Avoid complexities of event loops
- Avoid costs of real threads

Complexities of event loops

A simple

```
sock = socket.create_connection((host, 80))  
sock.sendall("GET / HTTP/1.0\r\n\r\n")  
data = sock.recv(1024)
```

becomes

Complexities of event loops

```
def connectionMade(self):
```

```
    ...
```

```
def dataReceived(self, data):
```

```
    ...
```

```
def connectionError(self, error):
```

```
    ...
```

async vs. sync

exceptions for I/O errors

```
try:
```

```
    connect()
```

```
except IOError:
```

```
    ...
```

```
def connectionMade(self):
```

```
    ...
```

```
def connectionError(self, error):
```

```
    ...
```

async vs. sync

context managers

with open('log', 'w') as log:

io_operation()

log("connected")

io_operation()

log("sent")

explicit state machine

log object is closed by "with"

async vs. sync

synchronous programming model

```
handle_result(func(params))
```

```
d = Deferred()
```

```
func(d)
```

```
d.add_callback(handle_result)
```

```
d.add_errback(handle_error)
```

Giving up

- exception handling
- context managers
- synchronous programming style
- 3rdparty libraries

A subset of Python without batteries.

Generators / PEP-3156?

- Help *somewhat*
- But not enough:

with conn.cursor() as curs:

 curs.execute(SQL)

cursor() needs to do I/O in __exit__

- No compatibility with threading / 3rdparty libraries

Coroutines

- generators: non-stackful coroutine
 - Only yield to parent
 - Need special syntax
 - Only saves the top frame
- greenlet: stackful coroutine
 - Switching is just a function call
 - Switching to any target
 - Saves the whole stack, like a thread

Why?

- Avoid complexities of event loops
- Avoid costs of real threads

Threads vs green threads

- creation
 - `thread.start_new_thread`: 28usec
 - `gevent.spawn`: 5usec
 - `gevent.spawn_raw`: 1usec
- does not matter if used via pool

Threads vs green threads

memory

- threads: 8MB of stack by default
- greenlet: only allocate what's actually used
 - 350 bytes per Python frame
 - 10-15KB
- does not matter since the memory is virtual
 - limits number of threads on 32bit arch

Gevent server

```
def handle(socket, addr):  
    # read out the request  
    for line in socket.makefile():  
        if not line.strip():  
            break  
    # send the response  
    socket.sendall(HTTP_RESPONSE)  
    socket.close()  
  
from gevent.server import StreamServer  
StreamServer(':5000', handle).serve_forever()
```

Threaded server

```
queue = Queue()
```

```
def worker():  
    while True:  
        socket = queue.get()  
        handle(socket)
```

```
for _ in xrange(1000):  
    thread.start_new_thread(handle, ())
```

```
while True:  
    socket, addr = listener.accept()  
    queue.put(socket)
```

Benchmark

`ab -n 1000 -c 100 http://localhost:5000/`

- threaded: 7.1K RPS, latency 14ms
- gevent: 9.3k RPS, latency 11ms
- The threaded server can probably be improved

Mutex / context switch benchmark

```
def func(source, dest, finished):
```

```
    source_id = id(source)
```

```
    for _ in xrange(COUNT):
```

```
        source.acquire()
```

```
        dest.release()
```

```
    finished.release()
```

```
thread.start_new_thread(func, (sem1, sem2, a_finished))
```

```
thread.start_new_thread(func, (sem2, sem1, b_finished))
```

Threads vs green threads

- context switch
 - 2 threads switch to each other using 2 mutexes
- gevent threads: 15ns
- real threads: 60ns
 - 2 CPUs – GIL contention

to avoid “taskset 1 python ...”

Threads vs green threads

- gevent threads: 15ns
- real threads: 12ns

PyPy:

- gevent threads: 11ns
- real threads: 7ns
- requires warmup, with only 100 iterations:
115ns / 35ns

ACTUALLY, I'M NOT EVEN MAD



THAT'S AMAZING

JOIN THE REVOLUTION

Thousands of Threads and Blocking I/O

The old way to write Java Servers is New again
(and way better)

Paul Tyma

paul.tyma@gmail.com

paultyma.blogspot.com

<http://www.mailinator.com/tymaPaulMultithreaded.pdf>



MARCH 3-7, 2008, SANTA CLARA, CA

Threads on Linux

- Threads used to have awful performance
- But since linux 2.6 / NPTL they've improved a lot
 - idle threads cost is almost zero
 - context switching is much faster
 - you can create lots of them
- (in Java) sync I/O is 30% faster than async I/O

cooperative OS threads?

- Already exclusive due to GIL
- `sys.setswitchinterval(2 ** 31)` # Python 3
- `sys.setcheckinterval(2 ** 30)` # Python 2

would not work if you have CPU-hungry threads, but neither will async frameworks

Can gevent be speed up?

- Switching and primitives are in Python
 - Let's try C
- Switching is done through Hub
 - Let's try directly

libgevent

- stacklet: a C module that PyPy uses to implement greenlet
- libuv: Node.JS event loop

libevent

```
gevent_cothread t1;  
gevent_cothread_init(hub, &t1, sleep0);  
gevent_cothread_spawn(&t1);  
gevent_wait(hub)
```


libgevent

- It's only a prototype
- `spawn()`, `sleep()`, `wait()`
- channels, semaphores
- `getaddrinfo()`
- naive Python wrapper

<https://github.com/denik/libgevent>

Threads vs green threads

- context switch
 - 2 threads switch to each other using 2 semaphores
- gevent threads: 15ns
- real threads: 12ns
- libgevent/gevent2: 1.8ns



Conclusions

- Thread pool is a deployment option with considering
- Gevent's performance can be pushed further
- Avoid framework lock-in
 - Migrating between gevent and threads is easy
 - Migrating between async and sync models is not
- The better threads are, the more irrelevant async frameworks are (gevent included)
 - And threads are already pretty fast

References

gevent:

<http://gevent.org>

faster gevent experiment:

<https://github.com/denik/libgevent>

pre-fork server for gevent:

<https://github.com/surfly/geventserver>

“Thousands of threads and blocking I/O” by Paul Tyma:

<http://www.mailinator.com/tymaPaulMultithreaded.pdf>